

Cues

Cues: the trustworthy AI shortcut layer for Android

A shortcut is a promise we can prove

Cues makes the path from a human request to a phone action visible. The user sees what the system understood, what the parser accepted, and what the runtime actually did.

Voice is the shortest path to intent

People already describe a desired outcome in natural language. Cues accepts speech or text in the user's language, then creates one canonical English representation for the parser without making language a new setup task.

Automation is powerful, but not trustworthy by default

Opaque automation creates three failure modes: the request is misunderstood, the device context is wrong, or the user cannot reconstruct what happened. Cues treats each failure as a state to surface, not a reason to hide the result. A signal that cannot be read stays unknown through evaluation and into the receipt — unknown never quietly reads as "no", and it never grants permission to act either.

Cues separates language from authority

Sarvam handles translation only. Gemma handles intent normalization and repair. The deterministic grammar parser remains the authority for executable syntax. A user must still review, approve, and arm the routine — explained before it acts, granted the least access it needs, and never running further than the user put it.

From spoken intent to an armed routine

The loop is deliberately staged: capture intent, translate to English, normalize with Gemma, parse into a cue, resolve device and time context, present a review, then arm and execute. Receipts record the outcome for Insights. Worked example: "When I arrive at Office, text Mum I'm home" becomes a place-arrival trigger and a drafted message — reviewed and handed to the user's own messaging app, never sent on the app's own authority.

```
app/src/main/kotlin/com/cues/app/ui/ask/AskScreen.kt
app/src/main/kotlin/com/cues/app/drafting/OnDeviceLlmDrafter.kt
app/src/main/kotlin/com/cues/app/drafting/SarvamChatDrafter.kt
app/src/main/kotlin/com/cues/app/net/SarvamClient.kt
core/src/main/kotlin/com/cues/core/drafting/GrammarParser.kt
app/src/main/kotlin/com/cues/app/runtime/PowerReceiver.kt
app/src/main/kotlin/com/cues/app/runtime/TimeTriggerReceiver.kt
app/src/main/kotlin/com/cues/app/ui/insights/InsightsScreen.kt
```

Why Cues wins the trust layer

Local-first: inference and decisions run on-device by default. Explain-before-acting: every routine renders as WHEN / IF / DO / UNTIL / RESTORE before it can arm. Least privilege: each permission request is derived from what the routine does, asked for once it creates visible value. Human control: nothing executes until the user approves and arms it. Provenance: every drafted routine names the drafter that produced it. Fail safely: a blocked action is a PARTIAL session, never a quiet success, and cleanup only ever releases what Cues itself owns. That combination turns a clever suggestion into a shortcut a user can safely keep.

Ship the loop, then expand the surface

The next milestone is a reliable daily loop across time, charger, and device context. After that foundation is stable, Cues can add richer sensors and more capable actions without weakening the safety boundary.